

part1-submission

July 6, 2024

1 MSA 2024 Phase 2 - Part 1

1.1 Introduction

The main goal of Part 1 is to load and preprocess the **W Store Sales** dataset and perform preliminary exploratory data analysis (EDA). This will involve cleaning the data, handling missing values, and gaining initial insights into the dataset through visualizations and summary statistics.

1.2 1. Import libraries and pre-define functions

```
[109]: %matplotlib inline

import sklearn
import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
import seaborn as sns
import matplotlib.pyplot as plt
from typing import List, Optional
import plotly.express as px
from plotly.subplots import make_subplots

def _to_human_readable(text:str):
    """
    Converts a label into a human readable form
    """
    return text.replace("_", " ")

def _prepare_labels(df:pd.DataFrame, labels:List[Optional[str]], replace_nones:
    ↪ bool=True):
    """
    Ensures labels are human readable.
    Automatically picks data if labels not provided explicitly
    """

    human_readable = {}

    if isinstance(replace_nones, bool):
```

```

        replace_nones = [replace_nones] * len(labels)

    for i in range(len(labels)):
        lab = labels[i]
        if replace_nones[i] and (lab is None):
            lab = df.columns[i]
            labels[i] = lab

        # make human-readable
        if lab is not None:
            human_readable[lab] = _to_human_readable(lab)

    return labels, human_readable

def box_and_whisker(df: pd.DataFrame,
                    label_x: Optional[str] = None,
                    label_y: Optional[str] = None,
                    label_x2: Optional[str] = None,
                    title=None,
                    y_axis_range: Optional[list] = None,
                    row: int = 1,
                    col: int = 1,
                    fig=None):
    """
    Creates a box and whisker plot on the provided fig object.

    df: The data
    label_x: What to group by. Defaults to None
    label_y: What to plot on the y axis. Defaults to count of df.columns[0]
    label_x2: If provided, splits boxplots into 2+ per x value, each with its
    ↪own colour
    title: Plot title
    y_axis_range: Limits for the y-axis
    row: Row position in subplot
    col: Column position in subplot
    fig: The plotly figure object where the plot will be drawn
    """

    # Automatically pick columns if not specified
    selected_columns, axis_labels = _prepare_labels(df, [label_x, label_y, ↪
    ↪label_x2], replace_nones=[False, True, False])

    # Create the box plot
    box_fig = px.box(df,
                     x=selected_columns[0],
                     y=selected_columns[1],
                     color=label_x2,

```

```

        labels=axis_labels,
        title=title)

# Update the layout to limit the y-axis if range is specified
if y_axis_range is not None:
    box_fig.update_layout(yaxis=dict(range=y_axis_range))

for trace in box_fig.data:
    fig.add_trace(trace, row=row, col=col)

```

1.3 2. Find all variables and understand them

1.3.1 2.1 Loading data

```

[110]: # read the w store sales dataset
ds_sales = pd.read_csv('./dataset/store_sales.csv', encoding='latin-1')

# preview the first 10 rows of the dataset
ds_sales.head(10)

```

```

[110]:
  Row ID      Order ID  Order Date  Ship Date      Ship Mode  Customer ID \
0      1  CA-2016-152156  11/8/2016  11/11/2016    Second Class  CG-12520
1      2  CA-2016-152156  11/8/2016  11/11/2016    Second Class  CG-12520
2      4  US-2015-108966  10/11/2015  10/18/2015    Standard Class  SO-20335
3      6  CA-2014-115812   6/9/2014   6/14/2014    Standard Class  BH-11710
4     11  CA-2014-115812   6/9/2014   6/14/2014    Standard Class  BH-11710
5     24  US-2017-156909   7/16/2017   7/18/2017    Second Class  SF-20065
6     25  CA-2015-106320   9/25/2015   9/30/2015    Standard Class  EB-13870
7     28  US-2015-150630   9/17/2015   9/21/2015    Standard Class  TB-21520
8     30  US-2015-150630   9/17/2015   9/21/2015    Standard Class  TB-21520
9     37  CA-2016-117590  12/8/2016  12/10/2016     First Class  GH-14485

```

```

      Customer Name      Segment      Country      City ... \
0      Claire Gute    Consumer  United States    Henderson ...
1      Claire Gute    Consumer  United States    Henderson ...
2  Sean O'Donnell    Consumer  United States  Fort Lauderdale ...
3  Brosina Hoffman    Consumer  United States    Los Angeles ...
4  Brosina Hoffman    Consumer  United States    Los Angeles ...
5  Sandra Flanagan    Consumer  United States    Philadelphia ...
6      Emily Burns    Consumer  United States      Orem ...
7  Tracy Blumstein    Consumer  United States    Philadelphia ...
8  Tracy Blumstein    Consumer  United States    Philadelphia ...
9      Gene Hale    Corporate  United States    Richardson ...

```

```

  Postal Code  Region      Product ID      Category  Sub-Category \
0      42420    South  FUR-BO-10001798  Furniture    Bookcases
1      42420    South  FUR-CH-10000454  Furniture      Chairs

```

2	33311	South	FUR-TA-10000577	Furniture	Tables
3	90032	West	FUR-FU-10001487	Furniture	Furnishings
4	90032	West	FUR-TA-10001539	Furniture	Tables
5	19140	East	FUR-CH-10002774	Furniture	Chairs
6	84057	West	FUR-TA-10000577	Furniture	Tables
7	19140	East	FUR-BO-10004834	Furniture	Bookcases
8	19140	East	FUR-FU-10004848	Furniture	Furnishings
9	75080	Central	FUR-FU-10003664	Furniture	Furnishings

	Product Name	Sales	Quantity \
0	Bush Somerset Collection Bookcase	261.9600	2
1	Hon Deluxe Fabric Upholstered Stacking Chairs,...	731.9400	3
2	Bretford CR4500 Series Slim Rectangular Table	957.5775	5
3	Eldon Expressions Wood and Plastic Desk Access...	48.8600	7
4	Chromcraft Rectangular Conference Tables	1706.1840	9
5	Global Deluxe Stacking Chair, Gray	71.3720	2
6	Bretford CR4500 Series Slim Rectangular Table	1044.6300	3
7	Riverside Palais Royal Lawyers Bookcase, Royal...	3083.4300	7
8	Howard Miller 13-3/4" Diameter Brushed Chrome ...	124.2000	3
9	Electrix Architect's Clamp-On Swing Arm Lamp, ...	190.9200	5

	Discount	Profit
0	0.00	41.9136
1	0.00	219.5820
2	0.45	-383.0310
3	0.00	14.1694
4	0.20	85.3092
5	0.30	-1.0196
6	0.00	240.2649
7	0.50	-1665.0522
8	0.20	15.5250
9	0.60	-147.9630

[10 rows x 21 columns]

1.3.2 2.2 Exploring the columns

```
[111]: # view the columns in the dataset
ds_sales.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2121 entries, 0 to 2120
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Row ID          2121 non-null  int64
1   Order ID        2121 non-null  object
2   Order Date      2121 non-null  object
```

```

3   Ship Date      2121 non-null  object
4   Ship Mode      2121 non-null  object
5   Customer ID    2121 non-null  object
6   Customer Name  2121 non-null  object
7   Segment        2121 non-null  object
8   Country        2121 non-null  object
9   City           2121 non-null  object
10  State          2121 non-null  object
11  Postal Code    2121 non-null  int64
12  Region         2121 non-null  object
13  Product ID     2121 non-null  object
14  Category       2121 non-null  object
15  Sub-Category   2121 non-null  object
16  Product Name   2121 non-null  object
17  Sales          2121 non-null  float64
18  Quantity       2121 non-null  int64
19  Discount       2121 non-null  float64
20  Profit         2121 non-null  float64
dtypes: float64(3), int64(3), object(15)
memory usage: 348.1+ KB

```

```
[112]: ds_sales.describe(include='all')
```

```

[112]:
count      Row ID      Order ID  Order Date  Ship Date      Ship Mode  \
count      2121.000000      2121      2121      2121      2121
unique      NaN      1764      889      960      4
top      NaN  US-2015-129007  9/5/2016  12/6/2017  Standard Class
freq      NaN      4      10      10      1248
mean      5041.643564      NaN      NaN      NaN      NaN
std      2885.740258      NaN      NaN      NaN      NaN
min      1.000000      NaN      NaN      NaN      NaN
25%      2568.000000      NaN      NaN      NaN      NaN
50%      5145.000000      NaN      NaN      NaN      NaN
75%      7534.000000      NaN      NaN      NaN      NaN
max      9991.000000      NaN      NaN      NaN      NaN

      Customer ID  Customer Name  Segment      Country      City  ...  \
count      2121      2121      2121      2121      2121  ...
unique      707      707      3      1      371  ...
top      SV-20365  Seth Vernon  Consumer  United States  New York City  ...
freq      15      15      1113      2121      192  ...
mean      NaN      NaN      NaN      NaN      NaN  ...
std      NaN      NaN      NaN      NaN      NaN  ...
min      NaN      NaN      NaN      NaN      NaN  ...
25%      NaN      NaN      NaN      NaN      NaN  ...
50%      NaN      NaN      NaN      NaN      NaN  ...
75%      NaN      NaN      NaN      NaN      NaN  ...

```

max	NaN	NaN	NaN	NaN	NaN	NaN	...
-----	-----	-----	-----	-----	-----	-----	-----

	Postal Code	Region	Product ID	Category	Sub-Category	\
count	2121.000000	2121	2121	2121	2121	
unique	NaN	4	375	1	4	
top	NaN	West	FUR-FU-10004270	Furniture	Furnishings	
freq	NaN	707	16	2121	957	
mean	55726.556341	NaN	NaN	NaN	NaN	
std	32261.888225	NaN	NaN	NaN	NaN	
min	1040.000000	NaN	NaN	NaN	NaN	
25%	22801.000000	NaN	NaN	NaN	NaN	
50%	60505.000000	NaN	NaN	NaN	NaN	
75%	90032.000000	NaN	NaN	NaN	NaN	
max	99301.000000	NaN	NaN	NaN	NaN	

	Product Name	Sales	Quantity	Discount	\
count	2121	2121.000000	2121.000000	2121.000000	
unique	380	NaN	NaN	NaN	
top	KI Adjustable-Height Table	NaN	NaN	NaN	
freq	18	NaN	NaN	NaN	
mean	NaN	349.834887	3.785007	0.173923	
std	NaN	503.179145	2.251620	0.181547	
min	NaN	1.892000	1.000000	0.000000	
25%	NaN	47.040000	2.000000	0.000000	
50%	NaN	182.220000	3.000000	0.200000	
75%	NaN	435.168000	5.000000	0.300000	
max	NaN	4416.174000	14.000000	0.700000	

	Profit
count	2121.000000
unique	NaN
top	NaN
freq	NaN
mean	8.699327
std	136.049246
min	-1862.312400
25%	-12.849000
50%	7.774800
75%	33.726600
max	1013.127000

[11 rows x 21 columns]

1.3.3 2.3 Remove the redundant columnns

Remove those redundant identifiers and columns which do not help our analysis

```
[113]: # make a copy to keep the original dataset
ds_sales_modified = ds_sales.copy()

# drop the useless identifier columns, category and country columns have only
↳ one unique value and ship date is redundant to order date, so we drop them
↳ as well,
columns_to_remove = ['Row ID', 'Customer ID', 'Customer Name', 'Order ID',
↳ 'Postal Code', 'Product ID', 'Ship Date', 'Category', 'Country', 'Product
↳ Name']
ds_sales_modified.drop(columns=columns_to_remove, inplace=True)

ds_sales_modified.head()
```

```
[113]:
```

	Order Date	Ship Mode	Segment	City	State	Region	\
0	11/8/2016	Second Class	Consumer	Henderson	Kentucky	South	
1	11/8/2016	Second Class	Consumer	Henderson	Kentucky	South	
2	10/11/2015	Standard Class	Consumer	Fort Lauderdale	Florida	South	
3	6/9/2014	Standard Class	Consumer	Los Angeles	California	West	
4	6/9/2014	Standard Class	Consumer	Los Angeles	California	West	

	Sub-Category	Sales	Quantity	Discount	Profit
0	Bookcases	261.9600	2	0.00	41.9136
1	Chairs	731.9400	3	0.00	219.5820
2	Tables	957.5775	5	0.45	-383.0310
3	Furnishings	48.8600	7	0.00	14.1694
4	Tables	1706.1840	9	0.20	85.3092

1.3.4 2.4 Handle date and categorical features

- Date features: convert to date types and extract date elements such as month, day of weeks, quarters, etc
- Categorical features: encode them

```
[114]: # convert the date columns to datetime
ds_sales_modified['Order Date'] = pd.to_datetime(ds_sales_modified['Order_
↳ Date'])
ds_sales_modified['Day of Week'] = ds_sales_modified['Order Date'].dt.dayofweek
ds_sales_modified['Year'] = ds_sales_modified['Order Date'].dt.year
ds_sales_modified['Month'] = ds_sales_modified['Order Date'].dt.month
ds_sales_modified['Quarter'] = ds_sales_modified['Order Date'].dt.quarter
ds_sales_modified['Week_begin'] = ds_sales_modified['Order Date'].dt.
↳ to_period('W').apply(lambda r: r.start_time)
ds_sales_modified['Month_begin'] = ds_sales_modified['Order Date'].dt.
↳ to_period('M').apply(lambda r: r.start_time)
ds_sales_modified['Quarter_begin'] = ds_sales_modified['Order Date'].dt.
↳ to_period('Q').apply(lambda r: r.start_time)
```

```

ds_sales_modified['Year_begin'] = ds_sales_modified['Order Date'].dt.
    ↳to_period('Y').apply(lambda r: r.start_time)

# Create a dictionary to store the encoders for each column
encoders = {}

# List of columns to label encode, label encoding can avoid creating too many
    ↳columns
columns_to_encode = ['Segment', 'Ship Mode', 'City', 'State', 'Region',
    ↳'Sub-Category']

# Label encode each column and store the encoder
for col in columns_to_encode:
    encoder = LabelEncoder()
    ds_sales_modified[col] = encoder.fit_transform(ds_sales[col])
    encoders[col] = encoder

print(ds_sales_modified.info())
# Preview the modified dataset
ds_sales_modified.head()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2121 entries, 0 to 2120
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Order Date            2121 non-null  datetime64[ns]
1   Ship Mode             2121 non-null  int64
2   Segment               2121 non-null  int64
3   City                  2121 non-null  int64
4   State                 2121 non-null  int64
5   Region                2121 non-null  int64
6   Sub-Category          2121 non-null  int64
7   Sales                 2121 non-null  float64
8   Quantity              2121 non-null  int64
9   Discount              2121 non-null  float64
10  Profit                2121 non-null  float64
11  Day of Week           2121 non-null  int32
12  Year                  2121 non-null  int32
13  Month                 2121 non-null  int32
14  Quarter               2121 non-null  int32
15  Week_begin            2121 non-null  datetime64[ns]
16  Month_begin           2121 non-null  datetime64[ns]
17  Quarter_begin         2121 non-null  datetime64[ns]
18  Year_begin            2121 non-null  datetime64[ns]
dtypes: datetime64[ns](5), float64(3), int32(4), int64(7)

```


memory usage: 281.8 KB
None

```
[114]: Order Date  Ship Mode  Segment  City  State  Region  Sub-Category  \
0 2016-11-08      2      0   137   15      2      0
1 2016-11-08      2      0   137   15      2      1
2 2015-10-11      3      0   108    8      2      3
3 2014-06-09      3      0   184    3      3      2
4 2014-06-09      3      0   184    3      3      3

      Sales  Quantity  Discount   Profit  Day of Week  Year  Month  Quarter  \
0   261.9600         2     0.00   41.9136           1  2016    11         4
1   731.9400         3     0.00  219.5820           1  2016    11         4
2   957.5775         5     0.45 -383.0310           6  2015    10         4
3    48.8600         7     0.00   14.1694           0  2014     6         2
4  1706.1840         9     0.20   85.3092           0  2014     6         2

      Week_begin  Month_begin  Quarter_begin  Year_begin
0 2016-11-07  2016-11-01   2016-10-01 2016-01-01
1 2016-11-07  2016-11-01   2016-10-01 2016-01-01
2 2015-10-05  2015-10-01   2015-10-01 2015-01-01
3 2014-06-09  2014-06-01   2014-04-01 2014-01-01
4 2014-06-09  2014-06-01   2014-04-01 2014-01-01
```

1.4 3. Clean data

1.4.1 3.1 check any for null values

```
[115]: ds_sales_modified.isnull().sum()
```

```
[115]: Order Date      0
      Ship Mode      0
      Segment      0
      City          0
      State         0
      Region        0
      Sub-Category  0
      Sales         0
      Quantity      0
      Discount      0
      Profit        0
      Day of Week   0
      Year          0
      Month         0
      Quarter       0
      Week_begin    0
      Month_begin   0
      Quarter_begin 0
```

```
Year_begin      0
dtype: int64
```

1.4.2 3.2 check for any NA values

```
[116]: ds_sales_modified.isna().sum()
```

```
[116]: Order Date      0
      Ship Mode      0
      Segment       0
      City          0
      State         0
      Region        0
      Sub-Category  0
      Sales         0
      Quantity      0
      Discount      0
      Profit        0
      Day of Week   0
      Year          0
      Month         0
      Quarter       0
      Week_begin    0
      Month_begin   0
      Quarter_begin 0
      Year_begin    0
      dtype: int64
```

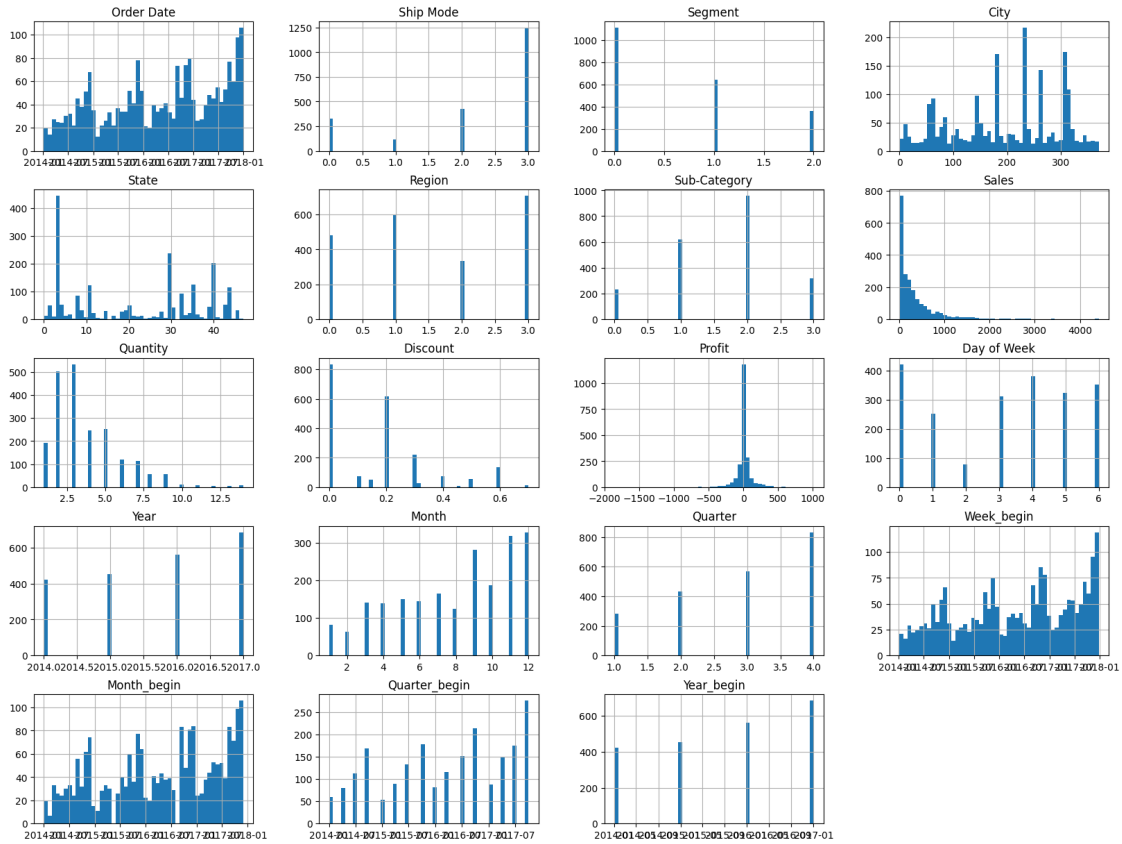
1.5 4. Visualise data and get some preliminary insights

1.5.1 4.1 distribution of all features

```
[117]: fig = ds_sales_modified.hist(bins=50, figsize=(20, 15))

      plt.suptitle('Chart 1, Distribution of all features', fontsize=16)
      plt.show()
```

Chart 1, Distribution of all features



The date related features indicate an increasing trend and season patterns. And city and state distributions highlight regional sales concentrations, with certain areas contributing more significantly.

1.5.2 4.2 Top sales by categorical features

```
[118]: ship_mode_data = ds_sales.groupby('Ship Mode')['Sales'].sum().
        ↪sort_values(ascending=True)
segment_data = ds_sales.groupby('Segment')['Sales'].sum().
        ↪sort_values(ascending=True)
city_data = ds_sales.groupby('City')['Sales'].sum().sort_values(ascending=True).
        ↪tail(10)
state_data = ds_sales.groupby('State')['Sales'].sum().
        ↪sort_values(ascending=True).tail(10)
region_data = ds_sales.groupby('Region')['Sales'].sum().
        ↪sort_values(ascending=True)
subcategory_data = ds_sales.groupby('Sub-Category')['Sales'].sum().
        ↪sort_values(ascending=True)
```

```

fig, axs = plt.subplots(2, 3, figsize=(24, 10))

fig.suptitle('Chart2. Sales Analysis by Various Categories', fontsize=16)

axs[0, 0].barh(ship_mode_data.index, ship_mode_data.values)
axs[0, 0].set_title('Sales by Ship Mode')

axs[0, 1].barh(segment_data.index, segment_data.values)
axs[0, 1].set_title('Sales by Segment')

axs[0, 2].barh(city_data.index, city_data.values)
axs[0, 2].set_title('Sales by City (Top 20)')

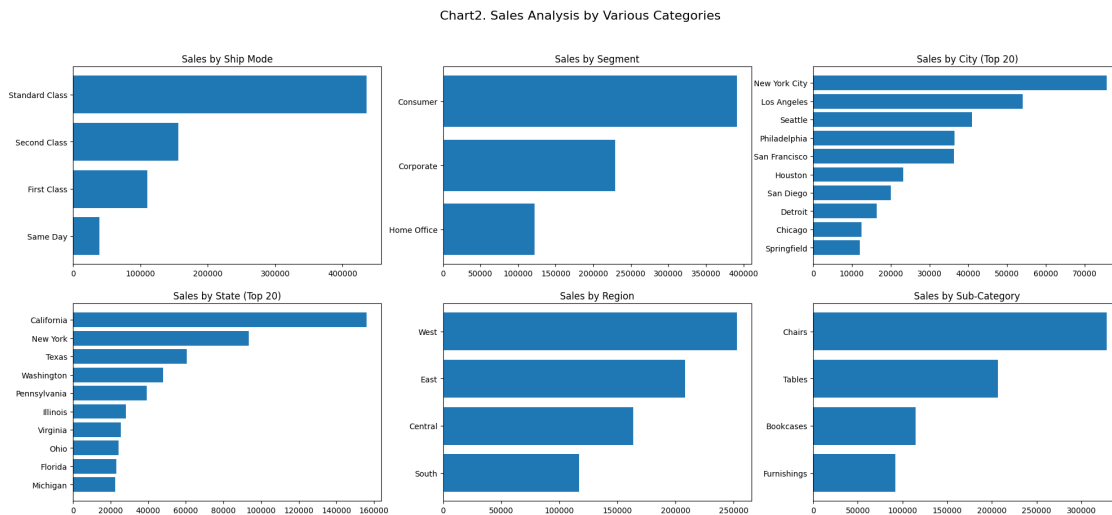
axs[1, 0].barh(state_data.index, state_data.values)
axs[1, 0].set_title('Sales by State (Top 20)')

axs[1, 1].barh(region_data.index, region_data.values)
axs[1, 1].set_title('Sales by Region')

axs[1, 2].barh(subcategory_data.index, subcategory_data.values)
axs[1, 2].set_title('Sales by Sub-Category')

plt.show()

```




Standard Class is the most common shipping method, and the Consumer segment leads in sales. New York City, Los Angeles, and Seattle are the top cities, while California, New York, and Texas are the top states. The West region outperforms other regions in sales. Among product sub-categories, Chairs generate the highest sales, followed by Tables and Bookcases.

1.5.3 4.3 Plot pair scatter charts and box charts

```
[119]: pairplot = sns.pairplot(ds_sales_modified, diag_kind='kde')
pairplot.fig.suptitle('Chart 3. Pairplot of all features', y=1.02, fontsize=20)

plt.show()
```



```
[120]: # Example usage
fig = make_subplots(rows=3, cols=3, subplot_titles=("Sales by Segment", "Sales_
↳by Ship Mode", "Sales by Region", "Sales by State", "Sales by City", "Sales_
↳by Sub-Category", "Sales by Quantity", "Sales by Discount"))
```

```

box_and_whisker(ds_sales_modified, label_x='Segment', label_y='Sales',
    ↳title='Sales by Segment', y_axis_range=[0, 1000], row=1, col=1, fig=fig)
box_and_whisker(ds_sales_modified, label_x='Ship Mode', label_y='Sales',
    ↳title='Sales by Ship Mode', y_axis_range=[0, 1000], row=1, col=2, fig=fig)
box_and_whisker(ds_sales_modified, label_x='Region', label_y='Sales',
    ↳title='Sales by Region', y_axis_range=[0, 1000], row=1, col=3, fig=fig)
box_and_whisker(ds_sales_modified, label_x='State', label_y='Sales',
    ↳title='Sales by State', y_axis_range=[0, 1000], row=2, col=1, fig=fig)
box_and_whisker(ds_sales_modified, label_x='City', label_y='Sales',
    ↳title='Sales by City', y_axis_range=[0, 1000], row=2, col=2, fig=fig)
box_and_whisker(ds_sales_modified, label_x='Sub-Category', label_y='Sales',
    ↳title='Sales by Sub-Category', y_axis_range=[0, 1000], row=2, col=3, fig=fig)
box_and_whisker(ds_sales_modified, label_x='Quantity', label_y='Sales',
    ↳title='Sales by Quantity', y_axis_range=[0, 1000], row=3, col=1, fig=fig)
box_and_whisker(ds_sales_modified, label_x='Discount', label_y='Sales',
    ↳title='Sales by Discount', y_axis_range=[0, 1000], row=3, col=2, fig=fig)

fig.update_layout(height=800, width=1200, title_text="Chart 4. Sales Analysis
    ↳by Various Categories")

fig.show()

```

```

[121]: # Get the original label for Sub-Category with encoded value 2
original_label = encoders['Sub-Category'].inverse_transform([2])[0]
print(f"The original label for Sub-Category with encoded value 2 is:
    ↳{original_label}")

```

The original label for Sub-Category with encoded value 2 is: Furnishings

Except the date patterns we've already found, from the box plots above we observed there's a sub-category['2']->'Furnishing', which has relative low sales.

1.6 4.4 Exploring time-series

```

[122]: # Aggregate Sales data by day of the week
weekly_sales_by_day = ds_sales_modified.groupby('Day of Week')['Sales'].sum()

# Aggregate Sales data by week
weekly_sales = ds_sales_modified.groupby('Week_begin')['Sales'].sum()

# Aggregate Sales data by month
monthly_sales = ds_sales_modified.groupby('Month_begin')['Sales'].sum()

# Aggregate Sales data by quarter
quarterly_sales = ds_sales_modified.groupby('Quarter_begin')['Sales'].sum()

# Aggregate Sales data by year

```

```

yearly_sales = ds_sales_modified.groupby('Year_begin')['Sales'].sum()

# Create a 2x3 grid of subplots with each subplot having a figsize of 10x5
fig, axs = plt.subplots(2, 3, figsize=(30, 10))

# Plot weekly sales by day of the week
axs[0, 0].bar(weekly_sales_by_day.index, weekly_sales_by_day, color='skyblue')
axs[0, 0].set_xlabel('Day of the Week')
axs[0, 0].set_ylabel('Sales')
axs[0, 0].set_title('Sales by Day of the Week')
axs[0, 0].set_xticks(weekly_sales_by_day.index)
axs[0, 0].set_xticklabels(weekly_sales_by_day.index, rotation=45)

# Plot weekly sales
axs[0, 1].plot(weekly_sales.index, weekly_sales, label='Weekly Sales',
              color='green')
axs[0, 1].set_xlabel('Date')
axs[0, 1].set_ylabel('Sales')
axs[0, 1].set_title('Weekly Sales')
axs[0, 1].legend()

# Plot monthly sales
axs[0, 2].plot(monthly_sales.index, monthly_sales, label='Monthly Sales',
              color='orange')
axs[0, 2].set_xlabel('Date')
axs[0, 2].set_ylabel('Sales')
axs[0, 2].set_title('Monthly Sales')
axs[0, 2].legend()

# Plot quarterly sales
axs[1, 0].plot(quarterly_sales.index, quarterly_sales, label='Quarterly Sales',
              color='red')
axs[1, 0].set_xlabel('Date')
axs[1, 0].set_ylabel('Sales')
axs[1, 0].set_title('Quarterly Sales')
axs[1, 0].legend()

# Plot yearly sales
axs[1, 1].plot(yearly_sales.index, yearly_sales, label='Yearly Sales',
              color='purple')
axs[1, 1].set_xlabel('Date')
axs[1, 1].set_ylabel('Sales')
axs[1, 1].set_title('Yearly Sales')
axs[1, 1].legend()

# Hide the empty subplot (bottom-right)

```


```

axs[1, 2].axis('off')

plt.suptitle('Chart 5. Sales Analysis by Time Periods', fontsize=20)

# Show the plots
plt.show()

```



We observed a strong seasonal pattern in the sales, so the next step is to plot the monthly autocorrelation and quarterly sub-series sales charts.

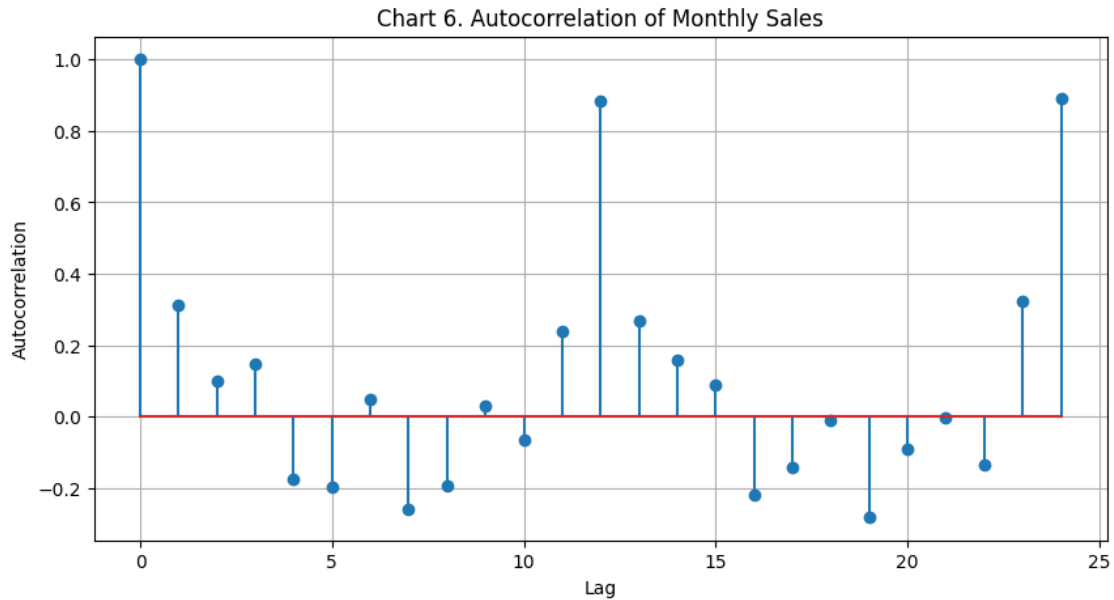
```

[123]: # Extracting monthly sales data
monthly_sales = ds_sales_modified.resample('ME', on='Order Date')['Sales'].sum()

# Calculate autocorrelation
lags = 24
autocorr = [monthly_sales.autocorr(lag) for lag in range(lags + 1)]

# Plotting autocorrelation
plt.figure(figsize=(10, 5))
plt.stem(range(lags + 1), autocorr)
plt.title('Chart 6. Autocorrelation of Monthly Sales')
plt.xlabel('Lag')
plt.ylabel('Autocorrelation')
plt.grid(True)
plt.show()

```

Autocorrelation result confirms the seasonal pattern, and there is a positive spike at lag 12.

```
[124]: from statsmodels.tsa.seasonal import seasonal_decompose

monthly_sales = ds_sales_modified.set_index('Order Date').
    ↪resample('ME')['Sales'].sum()

# Decompose the sales data
decomposition = seasonal_decompose(monthly_sales, model='additive')

# Extract the components
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

# Plot the components
plt.figure(figsize=(14, 10))
plt.title('Chart 7. Decomposition of Quarterly Sales')

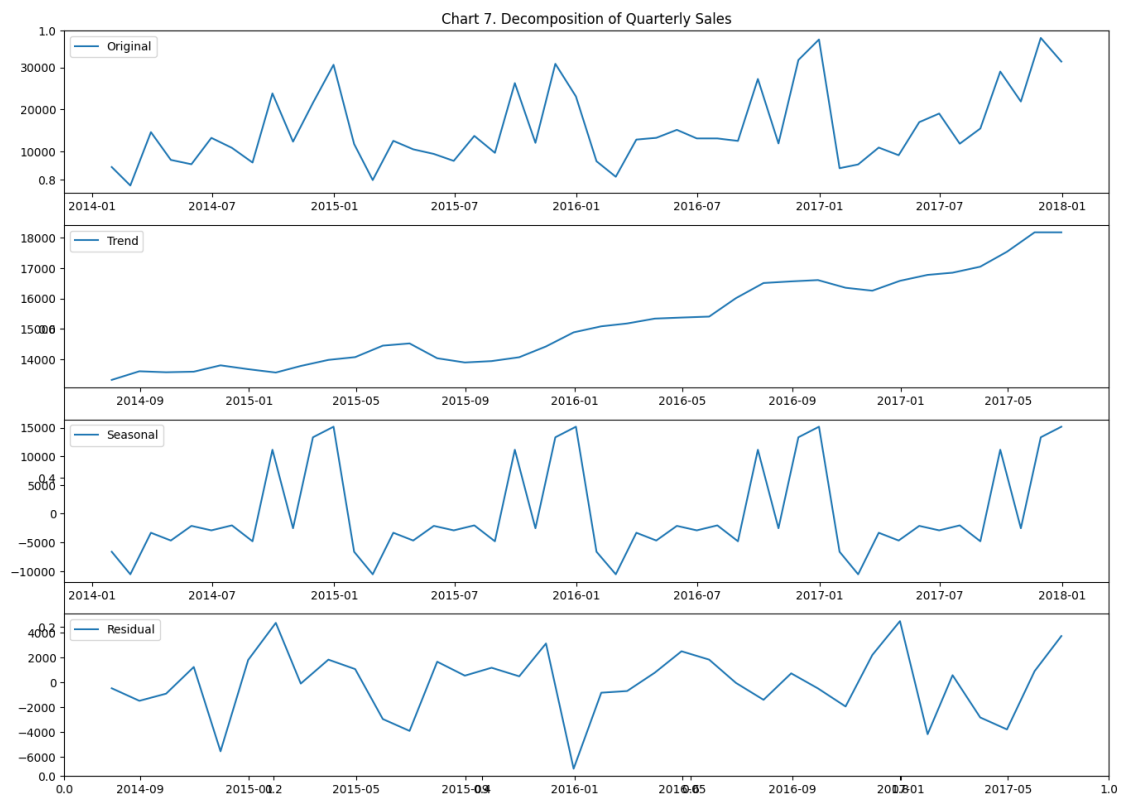
plt.subplot(411)
plt.plot(monthly_sales, label='Original')
plt.legend(loc='upper left')

plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='upper left')
```

```
plt.subplot(413)
plt.plot(seasonal, label='Seasonal')
plt.legend(loc='upper left')

plt.subplot(414)
plt.plot(residual, label='Residual')
plt.legend(loc='upper left')

plt.tight_layout()
plt.show()
```

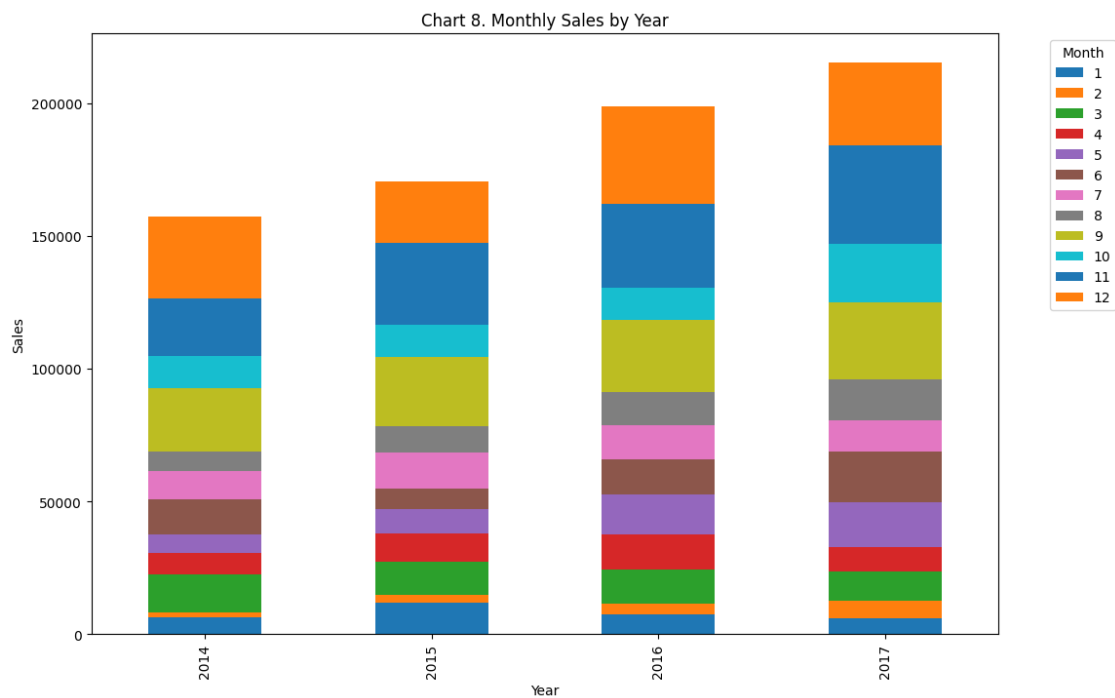


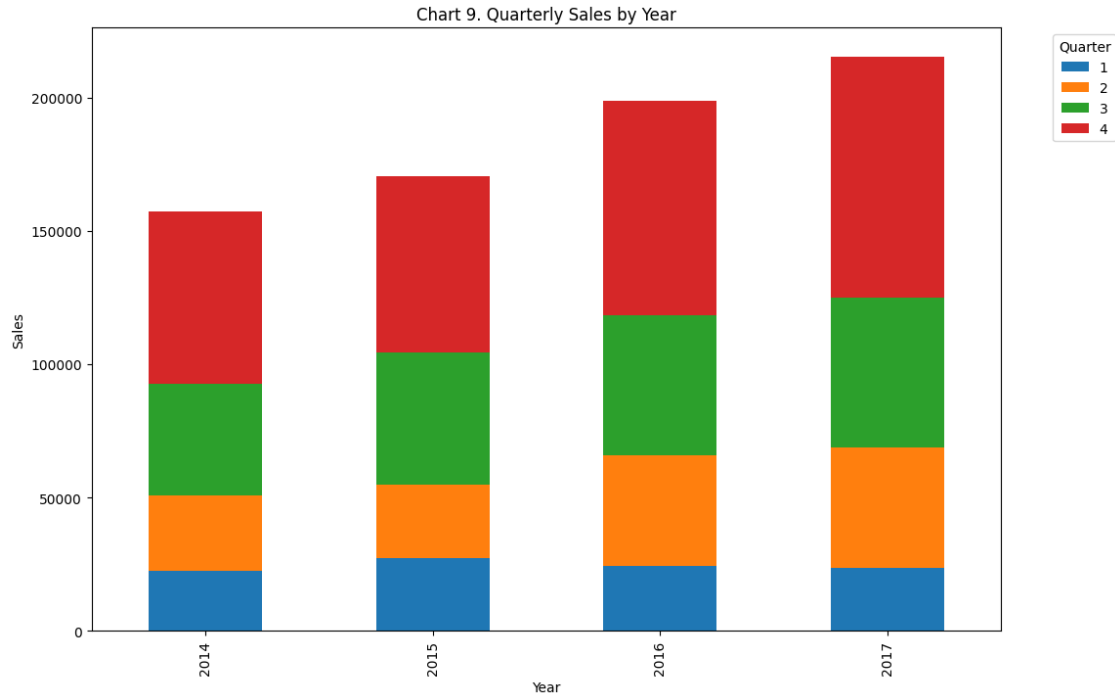
```
[125]: # Aggregate sales by month and year
monthly_sales = ds_sales_modified.groupby(['Year', 'Month'])['Sales'].sum().
    ↪unstack()

monthly_sales.plot(kind='bar', stacked=True, figsize=(12, 8))
plt.title('Chart 8. Monthly Sales by Year')
plt.xlabel('Year')
plt.ylabel('Sales')
plt.legend(title='Month', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```

```
# Aggregate sales by quarter and year
quarterly_sales = ds_sales_modified.groupby(['Year', 'Quarter'])['Sales'].sum().
    ↪unstack()

quarterly_sales.plot(kind='bar', stacked=True, figsize=(12, 8))
plt.title('Chart 9. Quarterly Sales by Year')
plt.xlabel('Year')
plt.ylabel('Sales')
plt.legend(title='Quarter', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```





The time-series plots revealed consistent sales growth, with noticeable peaks and troughs indicating seasonal variations. Further, we visualized the sales data using stacked bar charts, breaking down sales by month and quarter across different years. These visualizations highlighted the contributions of each month and quarter to the yearly sales, showing that the last quarter of each year consistently had higher sales, possibly due to holiday seasons.

1.7 5. Identify correlated variables

1.8 5.1 Feature selection

Based on the visulization results, we are going to drop following features: - Ship Date: It's related to the order date. We will focus on using the order date. - Ship Mode: It has no obvious relationship to the sales. - Customer ID, Row ID, Order ID, Product ID, Product Name, Customer Name: Identifiers that do not contribute to numerical analysis. - Region: It has no obvious relationship to the sales. - Segment: It has no obvious relationship to the sales. - Category: Single category presents in the dataset. - Country: Single country presents in the dataset. - Week/Month/Quarter/Year_begin: These columns are for plotting time-seris charts above. - Profit: This is derived from the sales.

```
[126]: # List of categorical features to be removed
categorical_features = ['Ship Mode', 'Segment', 'City', 'State', 'Region',
↳ 'Sub-Category', 'Day of Week', 'Year', 'Month', 'Quarter', 'Week_begin',
↳ 'Month_begin', 'Quarter_begin', 'Year_begin', 'Order Date']

# Remove categorical features from the DataFrame
```


```

ds_sales_numerical = ds_sales_modified.drop(columns=categorical_features)

# Calculate the correlation matrix for numerical features only
correlation_matrix_numerical = ds_sales_numerical.corr()

# Plot the heatmap
plt.figure(figsize=(14, 12))
sns.heatmap(correlation_matrix_numerical, annot=True, cmap='coolwarm',
            linewidths=0.5)
plt.title('Chart 10. Correlation Matrix of Numerical Features in W Store Sales_
Dataset')
plt.show()

```



From the correlation heatmap, we observed sales has relative strong relationship with profit and quantity, discount has strong relationship with region, city and profit. We are going to drop the

discount column.

From previous analysis, we found ship mode, segment, region, state (only use city), profit are less helpful for predicting sales, we are going to drop them as well.

```
[130]: ds_sales_selected = ds_sales_modified.drop(columns=['Discount', 'Ship Mode', 'Segment', 'Region', 'State', 'Profit'])
ds_sales_selected.to_csv('./dataset/store_sales_selected.csv', index=False)
ds_sales_selected.head()
```

```
[130]:   Order Date  City  Sub-Category    Sales  Quantity  Day of Week  Year  \
0  2016-11-08   137             0    261.9600         2           1  2016
1  2016-11-08   137             1    731.9400         3           1  2016
2  2015-10-11   108             3    957.5775         5           6  2015
3  2014-06-09   184             2     48.8600         7           0  2014
4  2014-06-09   184             3   1706.1840         9           0  2014
```

```
   Month  Quarter  Week_begin  Month_begin  Quarter_begin  Year_begin
0     11         4  2016-11-07  2016-11-01    2016-10-01  2016-01-01
1     11         4  2016-11-07  2016-11-01    2016-10-01  2016-01-01
2     10         4  2015-10-05  2015-10-01    2015-10-01  2015-01-01
3      6         2  2014-06-09  2014-06-01    2014-04-01  2014-01-01
4      6         2  2014-06-09  2014-06-01    2014-04-01  2014-01-01
```

1.9 6. Summary

1.9.1 Preprocess Steps

1. Data Loading and Initial Inspection:

- Loaded the dataset and displayed the first ten instances to understand the structure and content.
- Provided key statistical measures like mean and standard deviation.
- Encoded the categorical features.
- Visualized numerical columns through histograms to observe the distribution of values.

2. Data Cleaning:

- Checked for missing values and found none.

3. Visualization:

- Used pair plots to spot relationships between numerical features and sales, identifying meaningful relationships with quantity, city, and state.
- Plotted box charts to further confirm these relationships.
- Analyzed multi-item transactions but found nothing significant.
- Created new time-related columns like Day, Week, Month, Quarter, and Year from the Order Date column, revealing a strong seasonal pattern in sales data.

4. Correlation Analysis:

- Dropped irrelevant features and retained useful features based on the analysis results.

Through these steps, key insights into the dataset were gained, identifying important patterns and relationships, and preparing the data for further analysis or modeling.

1.9.2 Key Findings

- **Irrelevant features:** The following features were deemed not useful for predicting sales: 'Ship Date', 'Ship Mode', 'Customer ID', 'Customer Name', 'Region', 'Segment', 'Product Name', 'Profit', 'Row ID', 'Order ID', 'Country', 'Postal Code', 'Product ID', 'Category'.
- **Sales Patterns:** Most sales are less than \$50, with significant variations in sales amounts across different sub-categories.
- **Time Series Patterns:**
 - Sales increase over the years, showing a clear upward trend.
 - Higher sales on weekends, with Tuesday having the lowest sales.
 - More sales in September, November, and December, indicating a strong seasonal pattern.

1.9.3 Next Steps: Predicting Sales by Time and Selected Features